

Closest Pair in the Plane

Thore Husfeldt

September 12, 2016

Description

Find a closest pair of points in the plane using divide-and-conquer.

Requirements

As usual, your programme has to take its input either from standard input (STDIN) or from a specified file. It has to write to standard output (STDOUT).

Correctness Your solution must be consistent with the output files in the data directory. However, note that there is a possible source of confusion: Various programming languages support various data types for approximate real numbers. Typically, these types are called Double or Float. Depending on your choice of internal representation, the approximations will differ between implementations, so you might get a different output than me because of rounding. This cannot be helped; use your good judgement to decide if the error is based on rounding or on a programming mistake.¹

Your code has to be as clear and crisp as possible.

Your solution must run in time $O(n \log^2 n)$, where n is the number of points. If you want, you can implement a slightly more clever version that avoids sorting the input anew for every recursive call and runs in $O(n \log n)$ time. In fact, I found that version just as easy to write.

My solution runs in under a minute on all the inputs in total.

¹ This is a deliberately messy part of the exercise.

Tips

Have a look at the files – the format is easy enough to parse, but note that the positions of points in the plane is sometimes given as a floating-point number, and sometimes even in scientific notation.

Most programming languages have good support for this, e.g., Java's `Double.parseDouble` method is happy to process them. Also, note that the values are sometimes delimited by more than one space character. Again, Java's `String.split` class could be used to break them up. I like to do things by hand, so I used the following regex (in Perl):

```
$number = '[-+]?[d*\.]?[d+](?:[eE]([-+]?[d+])?';  
/((\d+)\s+($number)\s+($number))/;
```

It's probably a good idea to write the naïve quadratic-time algorithm first. It's nice to have access to The Truth for testing, and you need pretty much the exact same code at the bottom of the divide-and-conquer recursion anyway.

Deliverables

1. The source code for your implementation
2. A report in PDF. Use the report skeleton in the doc directory.